

Detecting patterns of speciation in the fossil records

pdf instructions

```
# Imports used in this project
from time import time # Checking time for optimiation
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import geopandas
import statsmodels.api as sm
world =
(geopandas.read_file(geopandas.datasets.get_path('naturalearth_lowres'
)))

# Constants for the project.
# Mammal Neogene (MN) time units. Doesn't include MQ18/MQ19 from the
table 1 of the pdf (like it shouldn't).
# Additional time units are pre-MN and post-MN
MN_UNITS = {'pre-MN': (1000000, 23), 'MN1': (23, 21.7), 'MN2': (21.7,
19.5), 'MN3': (19.5, 17.2), 'MN4': (17.2, 16.4), 'MN5': (16.4, 14.2),
'MN6': (14.2, 12.85), 'MN7-8': (12.85, 11.2), 'MN9':
(11.2, 9.9), 'MN10': (9.9, 8.9), 'MN11': (8.9, 7.6),
'MN12': (7.6, 7.1), 'MN13': (7.1, 5.3), 'MN14': (5.3,
5), 'MN15': (5, 3.55), 'MN16': (3.55, 2.5),
'MN17': (2.5, 1.9), 'post-MN': (1.9, 0)}

time_unit_order = MN_UNITS.keys() # Used to sort DataFrames for better
readability
```

Data preprocessing

1. First of all we need to get the data from the NOW database. Include the species list with comma as a field separator, and save the given output as `.txt` file to our project folder.
1. Create a pandas `DataFrame` that contains all of the data and save it as a `.csv` file. How many rows does the `DataFrame` contain?

```
def txt_to_csv():
    df = pd.read_csv(r'fossil_data.txt')
    df.to_csv(r'fossil_data.csv', index=None)
    return df

if __name__ == '__main__':
    df = txt_to_csv()
```

```
print(f'The DataFrame has: {len(df)} rows.')
print(df.head(2)) # Testing the DataFrame
```

The DataFrame has: 67693 rows.

	LIDNUM	NAME	LATSTR	LONGSTR	LAT	LONG
MAX_AGE \						
0	21390	Aarau	47 23 0 N	8 3 0 E	47.383	8.050000
14.2000						
1	27232	Aba Zawei	33 15 00 N	102 25 00 E	33.250	102.416667
0.0295						

	BFA_MAX	BFA_MAX_ABS	FRAC_MAX	...	MW_CS_ROUND	MW_CS_BLUNT	DIET_1	\
0	mn6	\N	\N	...	\N	\N	\N	
1	\N	C14	\N	...	\N	\N	p	

	DIET_2	DIET_3	LOCOM01	LOCOM02	LOCOM03	SPCOMMENT	SYNONYMS
0	\N	\N	\N	\N	\N	\N	\N
1	herbivore	graze	te	surficial	cursorial	\N	\N

[2 rows x 87 columns]

Answer

The DataFrame contains 67693 rows; excluding the column names.

Cleaning the data

1. This exercise cleans the data and makes it usable.

```
def calculate_mn(mi, ma):
    avg = np.mean([mi, ma]) # Calculate the mean age with MIN_AGE (mi)
    and the MAX_AGE (ma)
    for key, (ma, mi) in MN_UNITS.items():
        if mi <= avg < ma:
            return key
```

```
def identify_species(genus, species, SIN_DICT):
    g_s = ''.join((genus, species))
    UID = len(SIN_DICT)
    if g_s not in SIN_DICT:
        SIN_DICT[g_s] = UID
    return SIN_DICT[g_s]
```

```
def clean_df():
    df = pd.read_csv(r'fossil_data.csv')

    # Dropping rows based on instruction 3a
    df = df[(df.LAT != 0) & (df.LONG != 0)]
```

```

df = df[(df.SPECIES != 'sp.') & (df.SPECIES != 'indet.')]

# Adding MN column based on instruction 3b
df['MN'] = [calculate_mn(mi, ma) for mi, ma in zip(df['MIN_AGE'],
df['MAX_AGE'])]

# Editing time units based on instruction 3c. Only 'Can Llobateres
1' needs to be modified
df['MN'] = np.where(df['NAME'] == 'Can Llobateres 1', 'MN9',
df['MN'])

# Creating Species Identification Number (SIN) based on
instruction 3d
SIN_DICT = dict()
df['SIN'] = [identify_species(genus, species, SIN_DICT) for genus,
species in zip(df['GENUS'], df['SPECIES'])]

# Dropping rows with the same species at the same locality based
on instruction 3e
df = df.drop_duplicates(subset=['NAME', 'SIN'])
df.index = np.arange(0, len(df)) # Reset df index after removing
rows

return df

if __name__ == '__main__':
    df = clean_df()
    # Checking the amount of rows left, unique species and localities
    based on instruction 3f
    print(f'Rows left after step 3e: {len(df)}') # 49454 rows
    print(f'The amount of unique species: {max(df.SIN)+1}.') # 9849
    uniques. Index starts at 0.
    print(f'The amount of unique localities:
    {len(set(df.NAME.values))}.') # 5500 localities

Rows left after step 3e: 49454
The amount of unique species: 9849.
The amount of unique localities: 5500.

```

Answer

The DataFrame contains 49454 rows - excluding the column names -, 9849 unique species, and 5500 unique localities.

Notes: List comprehension is significantly faster than pandas own apply() method.

Occurrences

1. Creating a `DataFrame` that shows for each species how many occurrences it has in each time unit.

```
def build_MNs_for_species(idx, amt):  
    # Initialize a dictionary with default value as 0  
    columns = list(time_unit_order)  
    columns.extend(['SIN'])  
    occurrence_dict = dict.fromkeys(columns, 0)  
    sin, mn = idx  
    occurrence_dict['SIN'] = sin # Add Species ID  
    occurrence_dict[mn] = amt # Add amount of species found at certain  
    time period  
    return occurrence_dict  
  
def occurrences_in_MN():  
    main_df = clean_df()  
    MN_unique = main_df.groupby('SIN').MN.value_counts().to_frame() #  
    Count occurrence of every MN for each species  
    MN_occ = pd.DataFrame([build_MNs_for_species(idx, amt) for idx,  
    amt in zip(MN_unique.index, MN_unique.MN)]) # Build DataFrame  
    MN_occ = MN_occ.groupby('SIN', as_index=False).sum() # Merge the  
    same species together  
    return MN_occ  
  
if __name__ == '__main__':  
    df = occurrences_in_MN()  
    print(df.head(10)) # Testing DataFrame
```

	SIN	pre-MN	MN1	MN2	MN3	MN4	MN5	MN6	MN7-8	MN9	MN10	MN11
MN12 \												
0	0	0	0	0	0	0	0	1	0	0	0	0
0												
1	1	0	0	0	0	0	0	0	0	0	0	0
0												
2	2	0	0	0	0	0	0	0	0	0	0	0
0												
3	3	0	0	0	0	0	0	0	0	0	0	0
0												
4	4	0	0	0	0	0	0	0	0	0	0	0
0												
5	5	0	0	0	0	0	0	0	0	0	0	0
0												
6	6	0	0	0	0	0	0	0	0	0	0	0
0												
7	7	0	0	0	0	0	0	0	0	0	0	0
0												
8	8	0	0	0	0	2	1	12	26	36	68	10

13												
9	9	0	0	0	0	1	1	14	35	3	0	0
0												
	MN13	MN14	MN15	MN16	MN17	post-MN						
0	0	0	0	0	0	0						
1	0	0	0	0	0	20						
2	0	0	0	0	0	30						
3	0	0	0	0	0	4						
4	0	2	11	1	0	0						
5	1	5	19	7	4	4						
6	0	1	2	0	0	0						
7	1	1	3	0	0	0						
8	3	0	7	0	0	0						
9	0	0	0	0	0	0						

Creating a `DataFrame` that shows for each species the time unit when it is first observed

```
def species_first_occurences():
    all_occ = occurrences_in_MN().set_index('SIN')
    first_occ = pd.DataFrame()
    first_occ['MN'] = all_occ.ne(0).idxmax(1) # Find the first
    observation for every species
    # Add count column for easier use for the next task
    first_occ['COUNT'] = all_occ.replace(0, np.nan).bfill(1).iloc[:,
0].astype(int)
    return first_occ

def proportion_of_first_occurences():
    main_df = clean_df()
    species = species_first_occurences()

    time_units_first_occ = species['MN'].value_counts()
    [time_unit_order] # First observation for every species
    time_units_all_occ = main_df.MN.value_counts()[time_unit_order] #
    All occurences in a dataframe
    all_species = species.groupby('MN').COUNT.sum()[time_unit_order] #
    All first observations of species

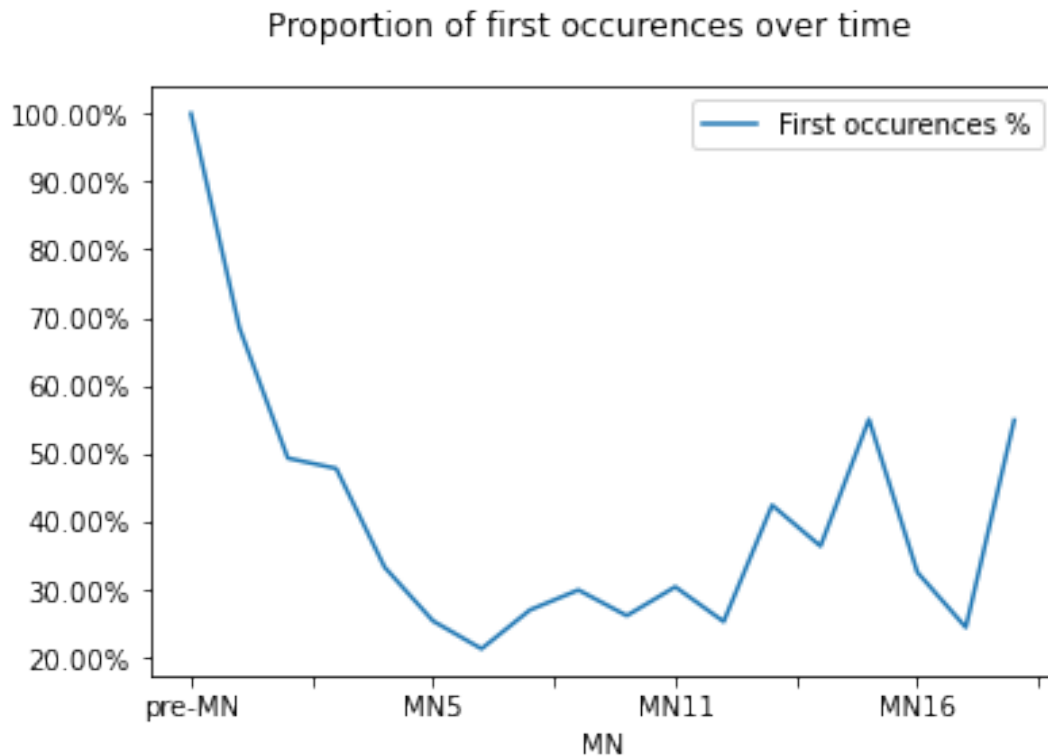
    proportion = pd.Series(all_species / time_units_all_occ)
    return proportion, time_units_all_occ # Return a tuple for easier
    plotting

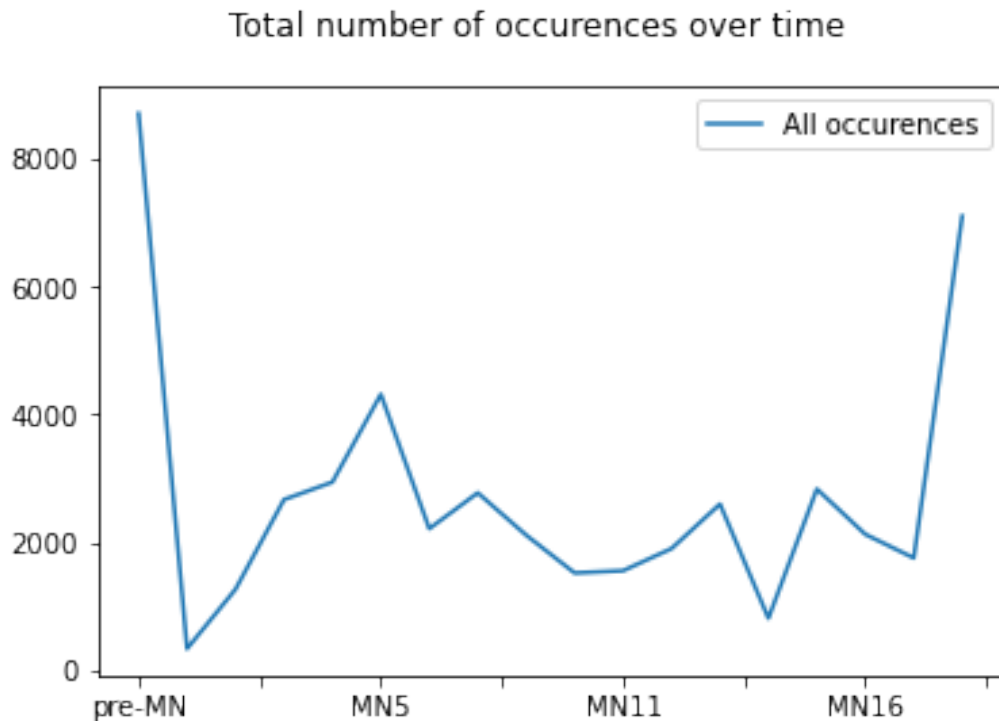
if __name__ == '__main__':
    proportion, all_occurences = proportion_of_first_occurences()
    plot1 = plt.figure(1)
    ax1 = proportion.plot()
    to_percentage = ax1.get_yticks()
    ax1.set_yticklabels([f'{val:.2%}' for val in to_percentage]) #
```

Format to percentage

```
plt.suptitle('Proportion of first occurrences over time')  
ax1.legend(['First occurrences %'])
```

```
plot1 = plt.figure(2)  
ax2 = all_occurences.plot()  
plt.suptitle('Total number of occurrences over time')  
ax2.legend(['All occurrences'])
```





Answer

There's really nothing to add to the plots. I think the wording of the assignment was a bit ambiguous for some parts. It can be understood in few different ways. I'm optimistic that I got it right though.

Notes: `time_unit_order` is a constant that was initialize on the top of the notebook. It helps to debug the code with better readability.

1. Creating a new `DataFrame` that collects the following information for every locality: locality number (LIDNUM), longitude, latitude, time unit, number of first occurrences in the locality, number of all occurrences in the locality and proportion of first occurrences in the locality. We are then going to plot some of the gathered information.

```
def location_df():
    main_df = clean_df()[['LIDNUM', 'LONG', 'LAT', 'MN', 'SIN']] #
    Taking SIN for easier DF building
    species = species_first_occurences().to_dict() # Assign to dict
    for easier use

    # Temporary column which defines if a row is a first occurence or
    not
    main_df['TEMP'] = [1 if species['MN'][sin] == mn else 0 for sin,
    mn in zip(main_df['SIN'], main_df['MN'])]
    main_df['FIRST_OCC'] = main_df.groupby(['LIDNUM'])
    ['TEMP'].transform('sum') # Amount of first occurrences in the locality
    main_df['ALL_OCC'] = main_df.groupby(['LIDNUM'])
    ['LIDNUM'].transform('count') # All occurrences of the locality
```

```

    main_df['PROPORTION'] = round((main_df['FIRST_OCC'] /
main_df['ALL_OCC']), 2)

    main_df.drop(columns=['TEMP', 'SIN'], inplace=True)
    main_df.drop_duplicates(subset=['LIDNUM'], inplace=True)
    return main_df

if __name__ == '__main__':
    df = location_df()
    print(df.head(20))

```

	LIDNUM	LONG	LAT	MN	FIRST_OCC	ALL_OCC
PROPORTION						
0	21390	8.050000	47.383000	MN6	1	1
1.00						
1	27232	102.416667	33.250000	post-MN	3	3
1.00						
4	27955	-1.544182	39.209991	MN15	0	4
0.00						
8	26550	46.533333	38.516667	MN12	0	1
0.00						
9	28578	1.788000	41.532600	MN7-8	2	9
0.22						
18	28579	1.788900	41.532600	MN7-8	1	10
0.10						
28	28580	1.788900	41.532600	MN7-8	2	4
0.50						
32	28581	1.788900	41.532600	MN7-8	3	8
0.38						
40	28582	1.788900	41.532600	MN7-8	5	17
0.29						
57	28583	1.788900	41.532600	MN7-8	6	16
0.38						
73	28584	1.780000	41.530000	MN7-8	6	11
0.55						
84	28585	1.788900	41.532600	MN7-8	3	6
0.50						
90	28586	1.788900	41.532600	MN7-8	4	9
0.44						
99	28587	1.788900	41.532600	MN7-8	2	14
0.14						
113	28244	1.765952	41.534164	MN7-8	1	1
1.00						
114	28245	1.765952	41.534164	MN7-8	1	1
1.00						
115	28246	1.765952	41.534164	MN7-8	1	1
1.00						
116	28260	1.679613	41.530754	post-MN	7	12
0.58						
128	27566	-0.558731	41.534907	MN2	0	1


```

0.00
129 28895 -69.000000 -17.300000 MN11 1 1
1.00

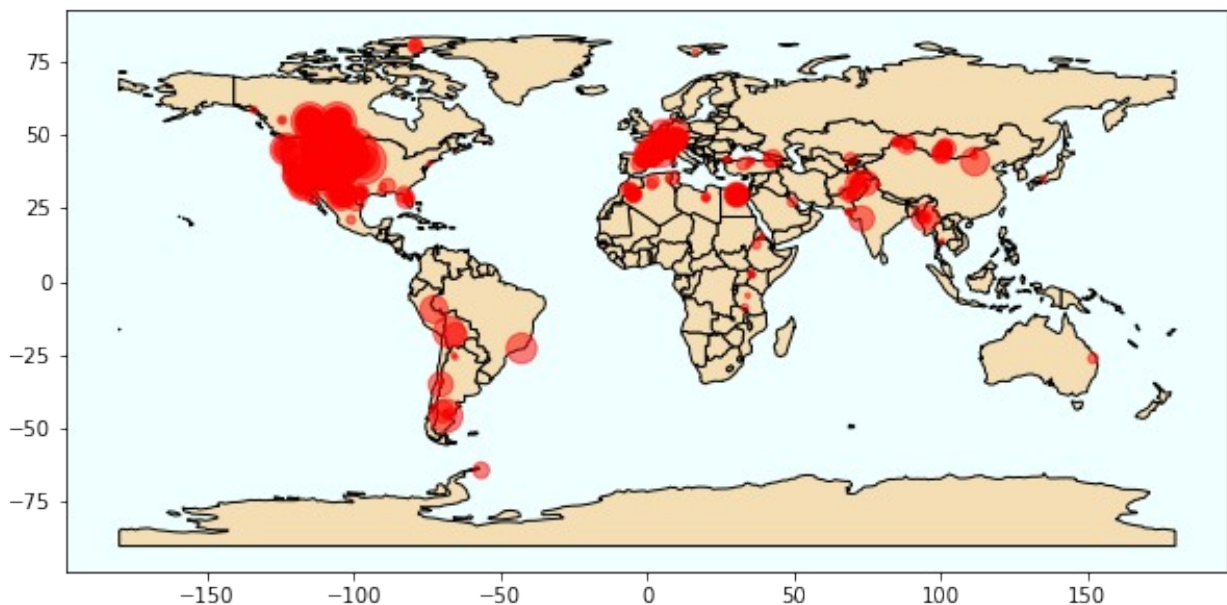
# Plotting localities in every time unit based on instruction 5 b.
def plot_locations():
    locations = location_df()
    for MN in MN_UNITS:
        fig, ax = plt.subplots(figsize=(10, 5))
        world.plot(ax=ax, color='wheat', edgecolor='black')
        ax.set_facecolor('azure')

        time_unit = locations[locations['MN'] == MN] # Get localities
        for specific time unit
            plt.suptitle(f'Time unit: {MN}. Localities: {len(time_unit)}')
            # Plot coordinates with dynamic marker size
            plt.scatter(x=time_unit.LONG, y=time_unit.LAT,
s=time_unit.ALL_OCC*6, c='red', alpha=0.5)
            plt.show()

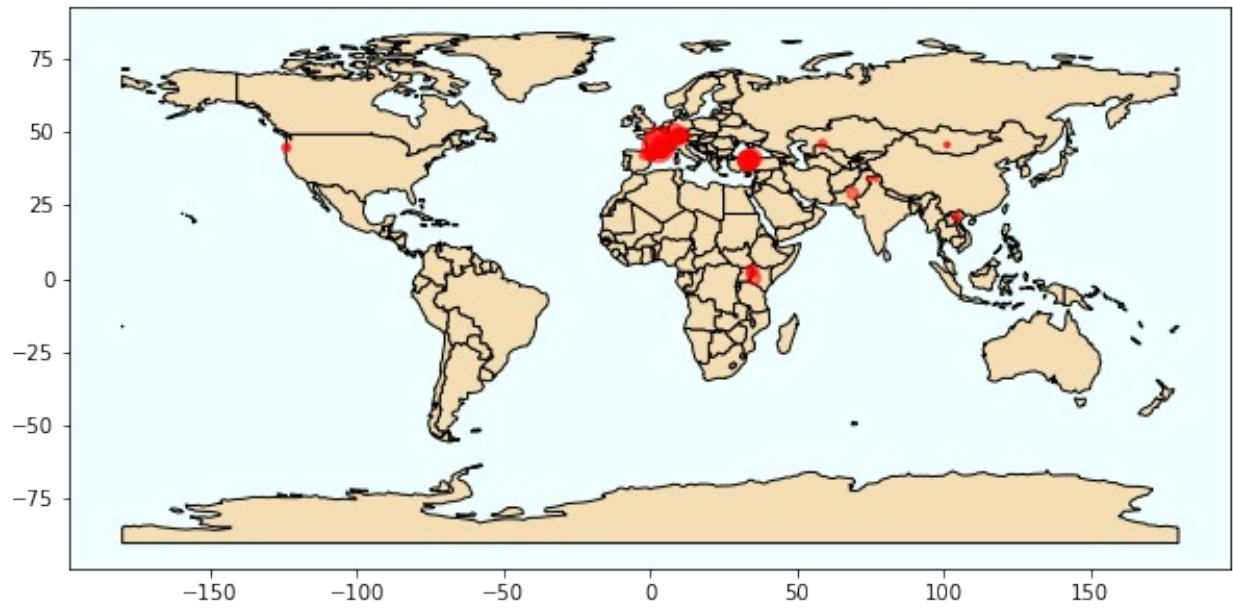
if __name__ == '__main__':
    plot_locations()

```

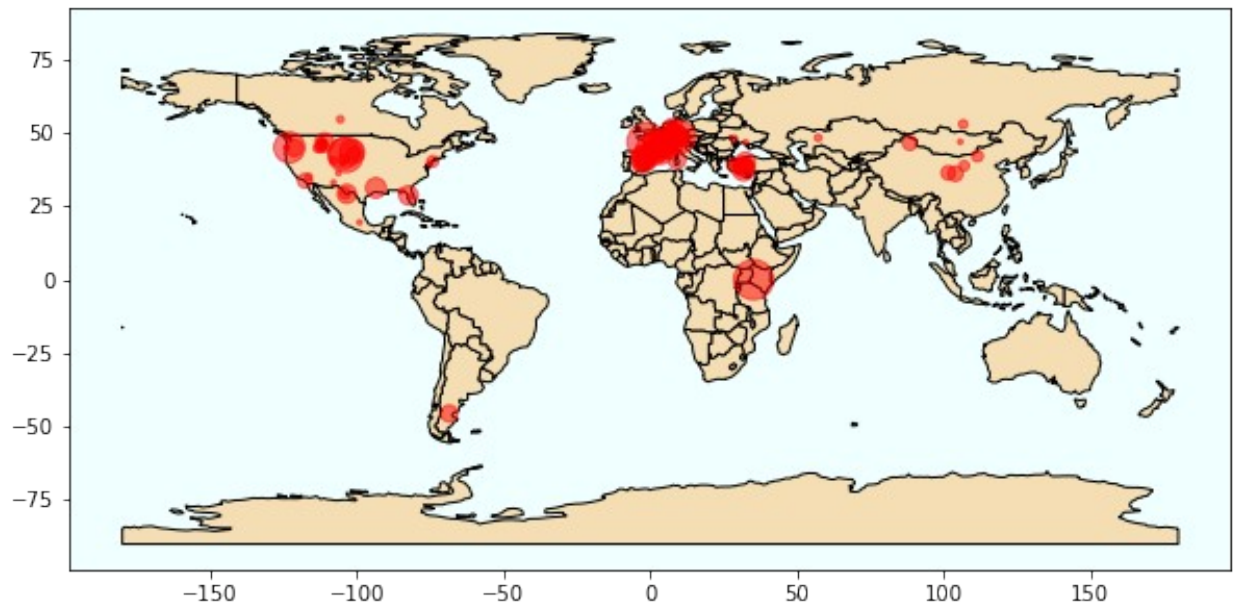
Time unit: pre-MN. Localities: 721



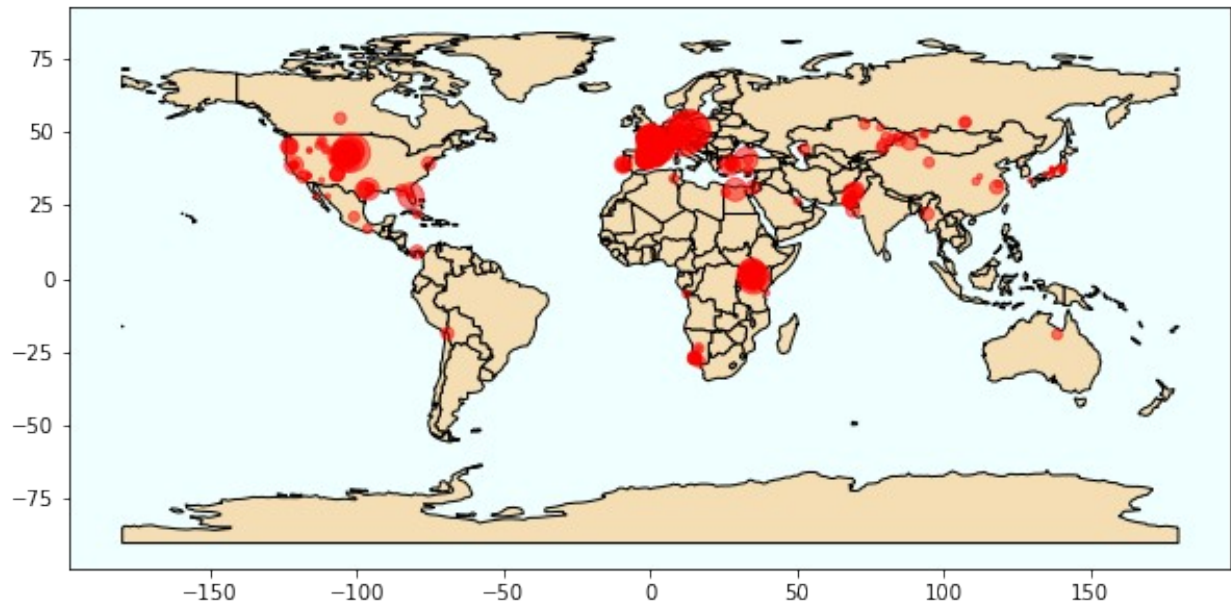
Time unit: MN1. Localities: 52



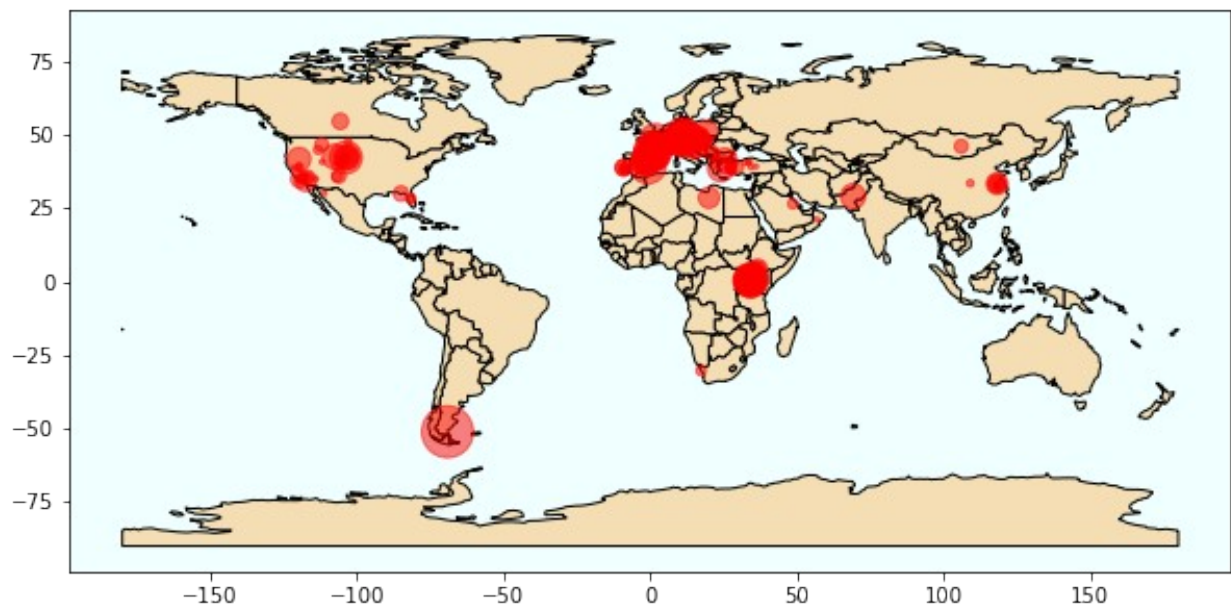
Time unit: MN2. Localities: 164



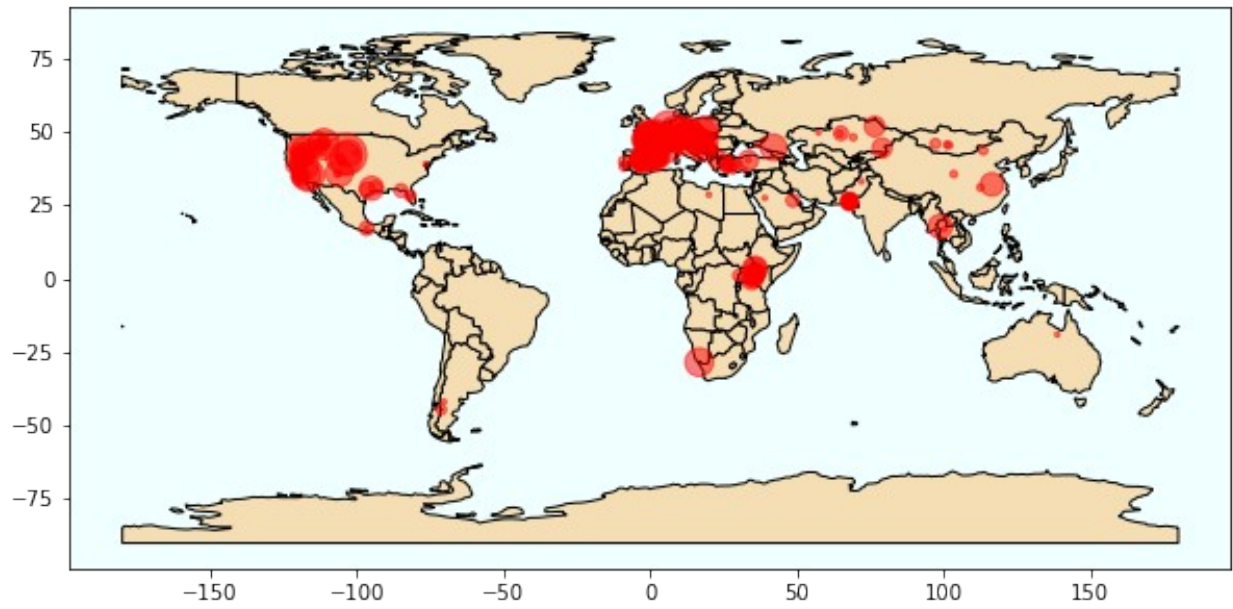
Time unit: MN3. Localities: 334



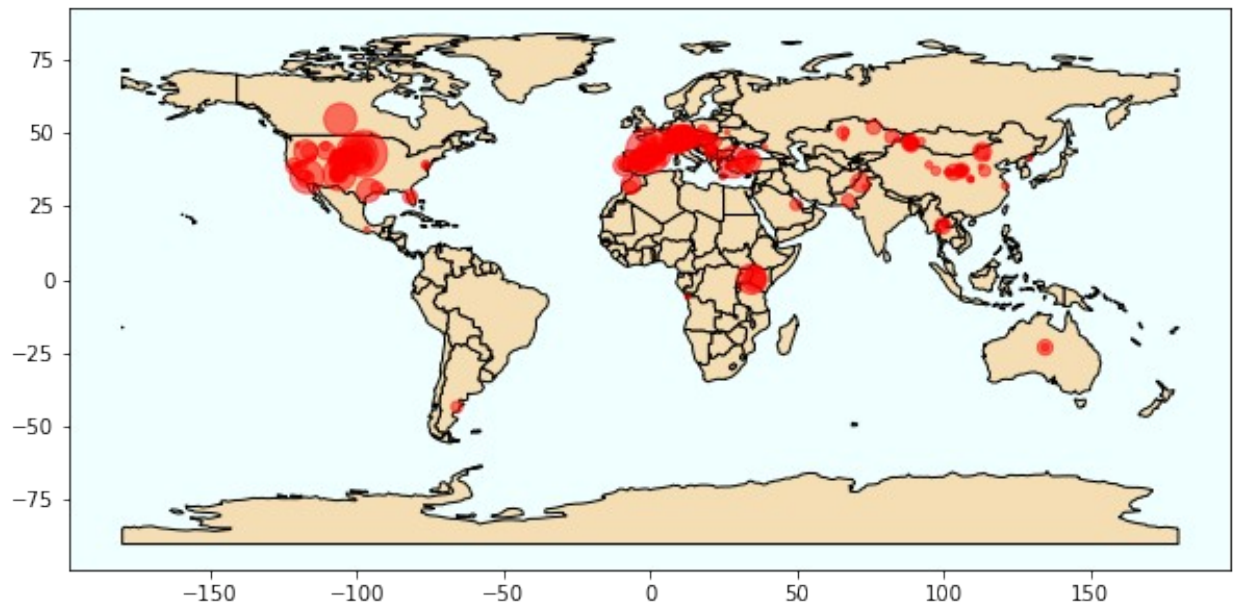
Time unit: MN4. Localities: 300



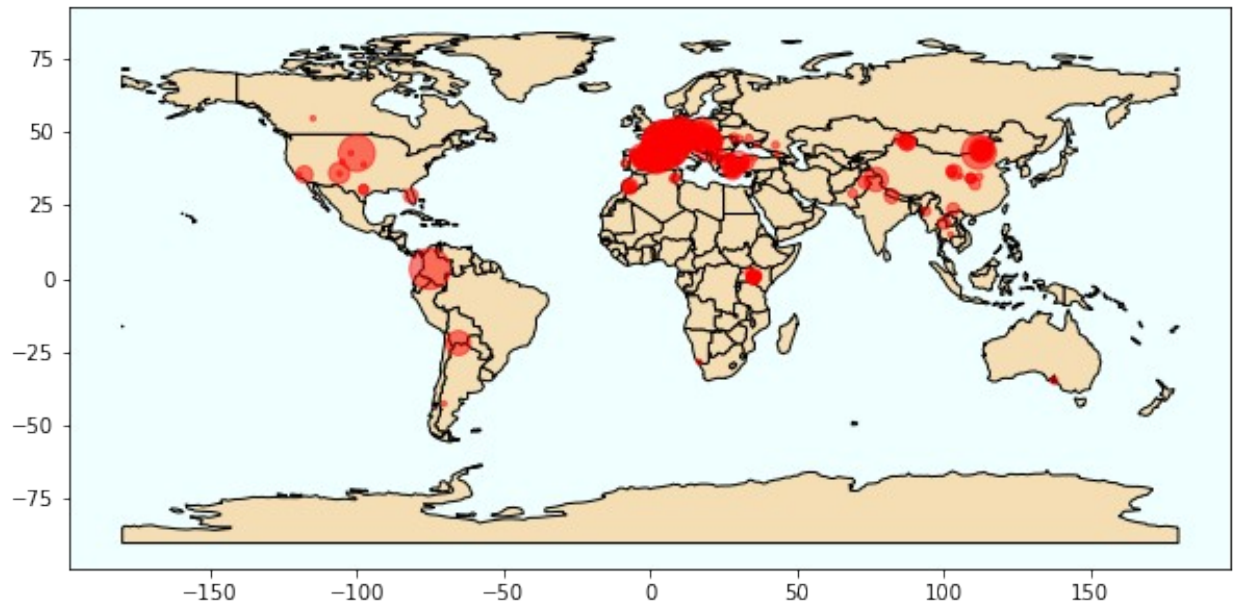
Time unit: MN5. Localities: 439



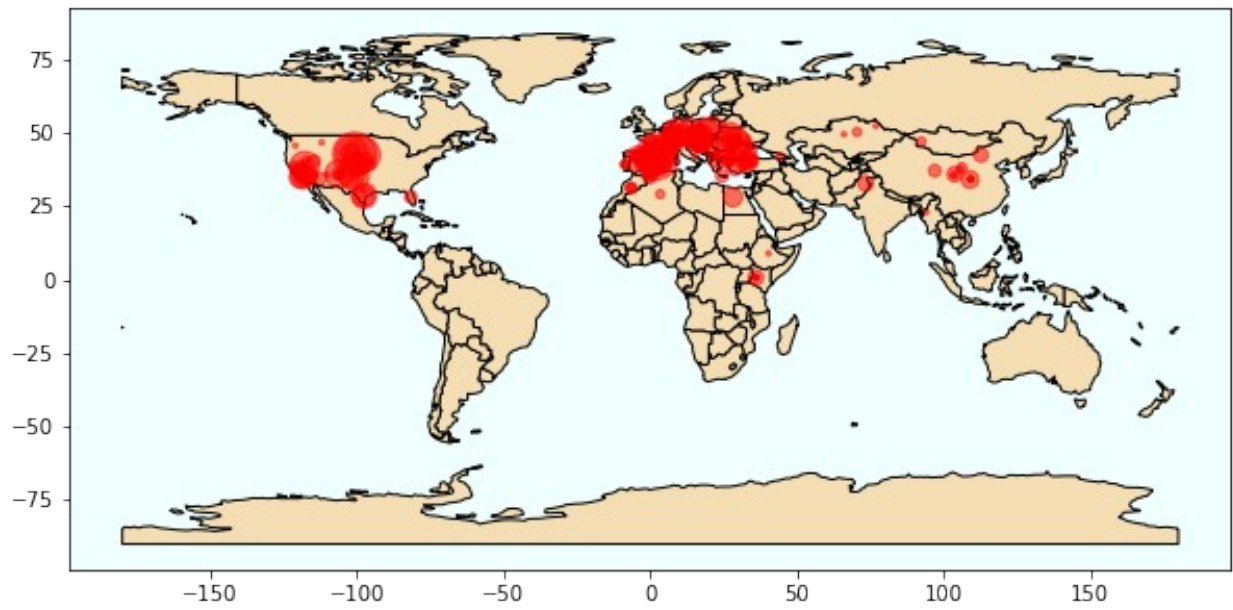
Time unit: MN6. Localities: 290



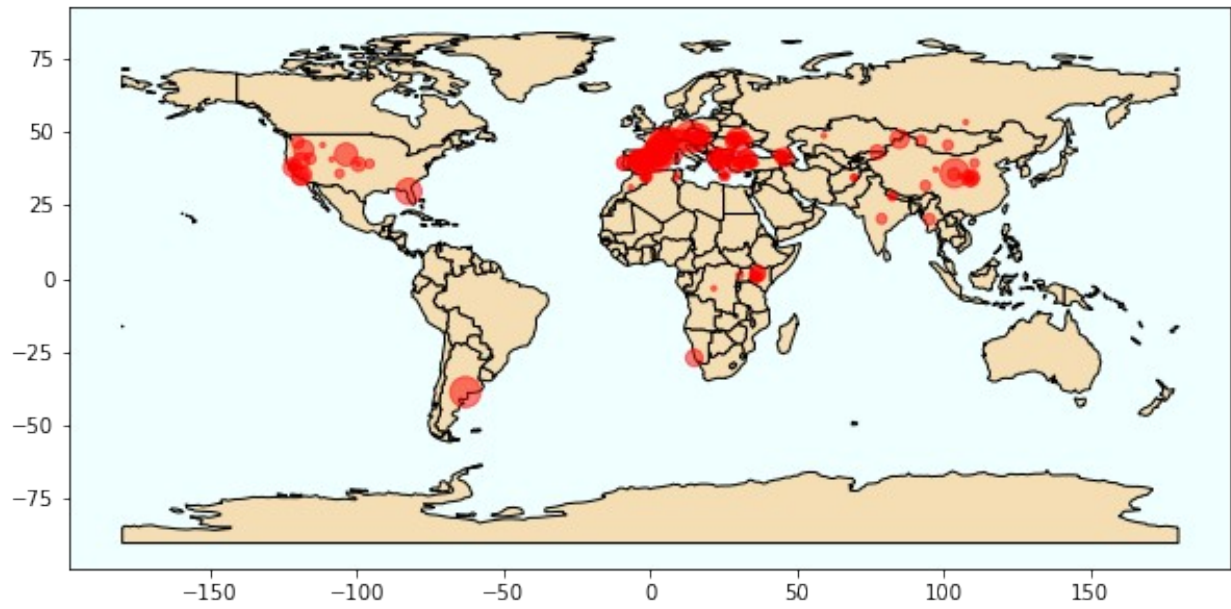
Time unit: MN7-8. Localities: 322



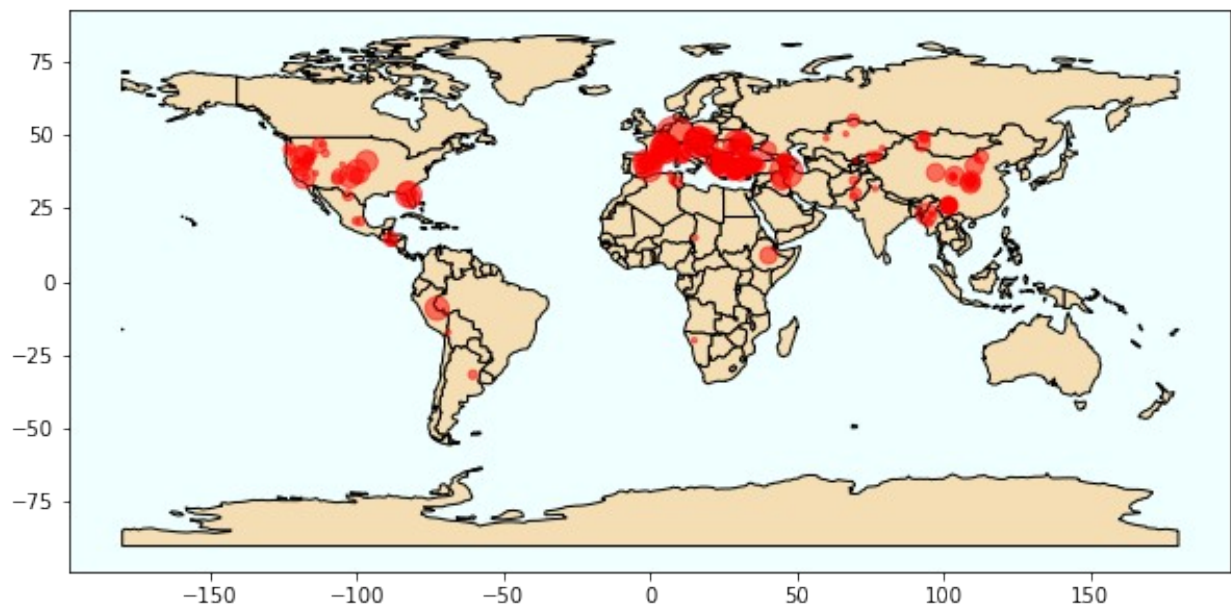
Time unit: MN9. Localities: 273



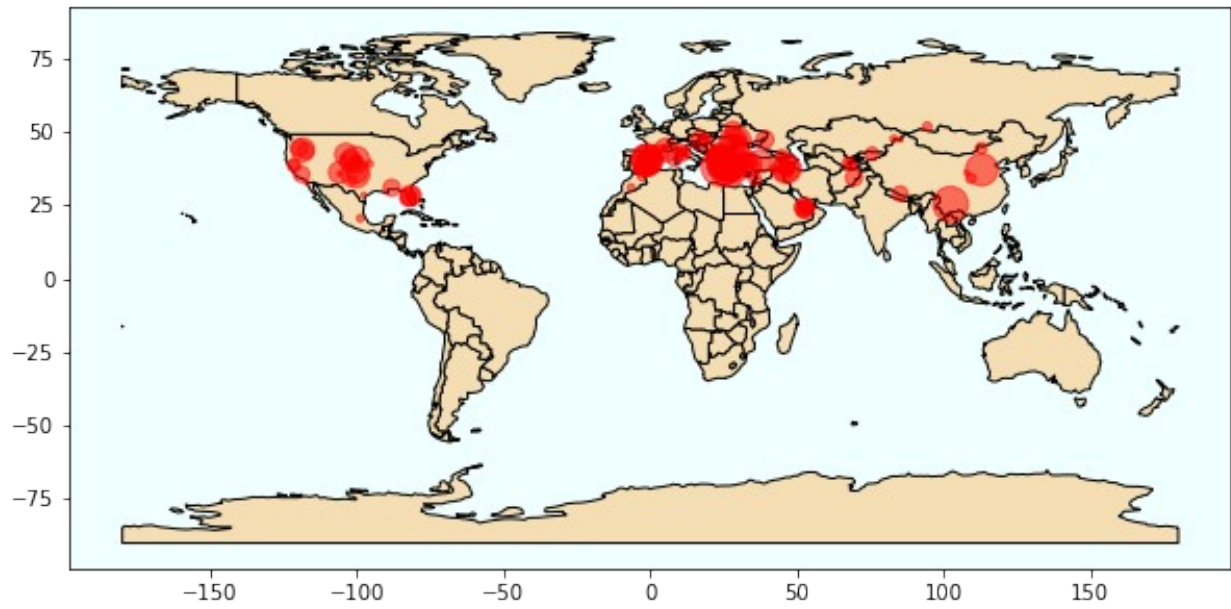
Time unit: MN10. Localities: 275



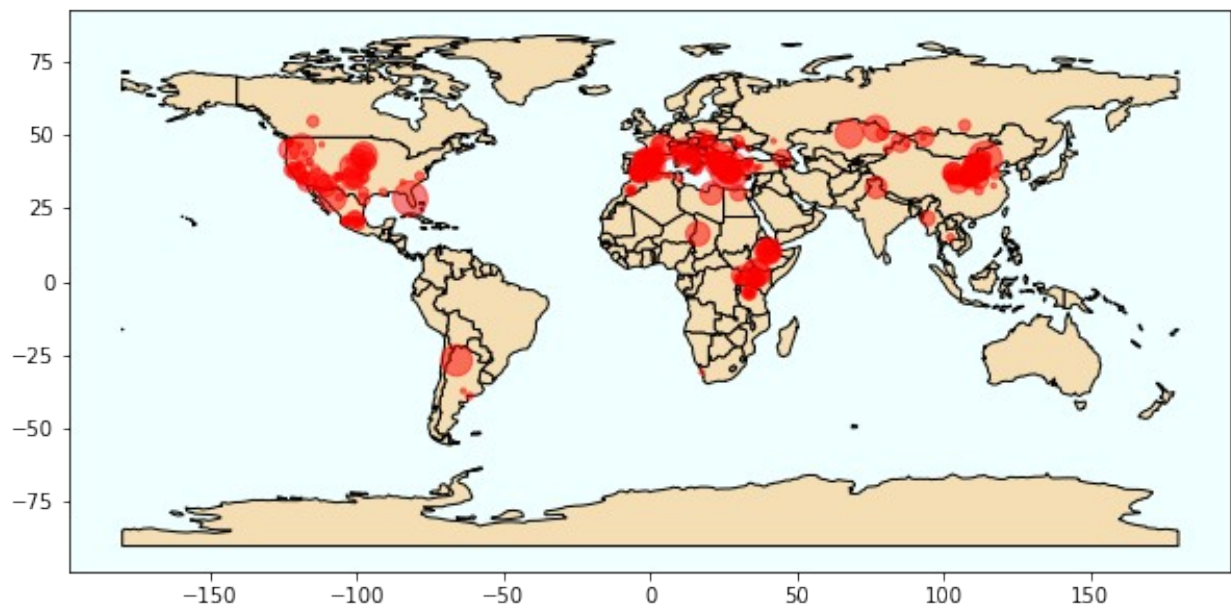
Time unit: MN11. Localities: 235



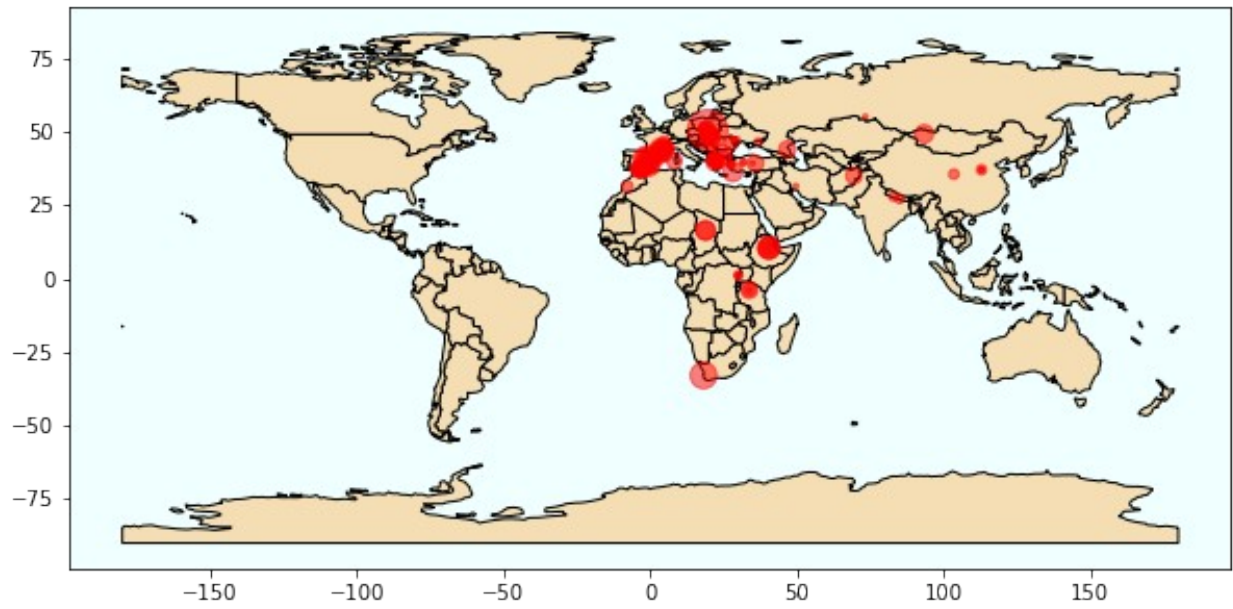
Time unit: MN12. Localities: 212



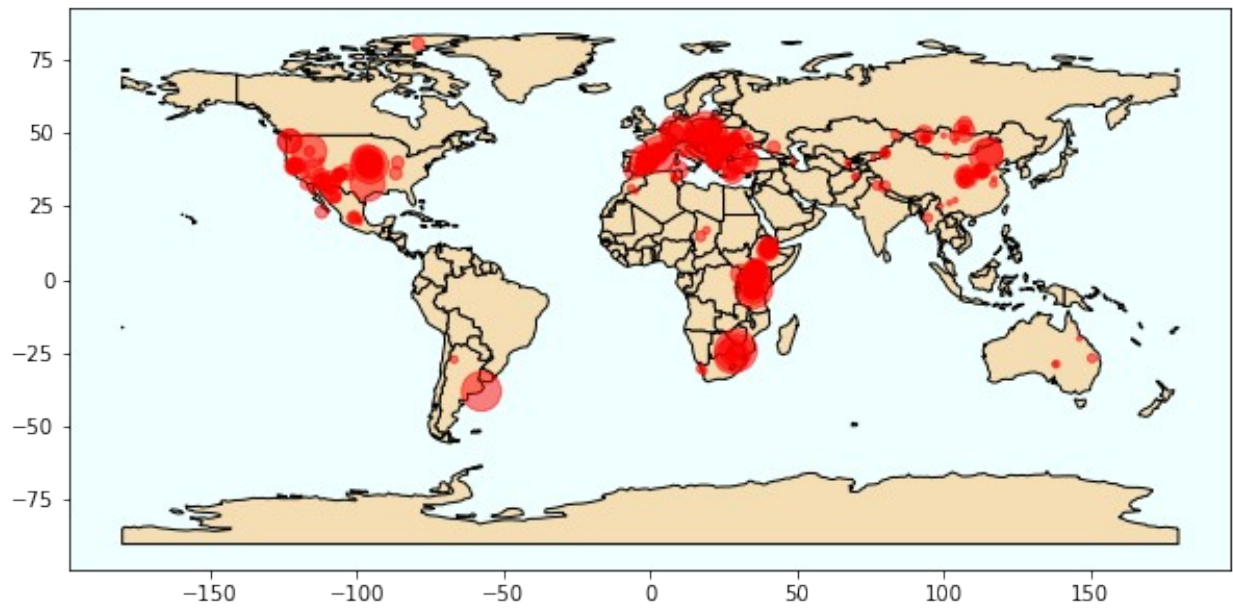
Time unit: MN13. Localities: 339



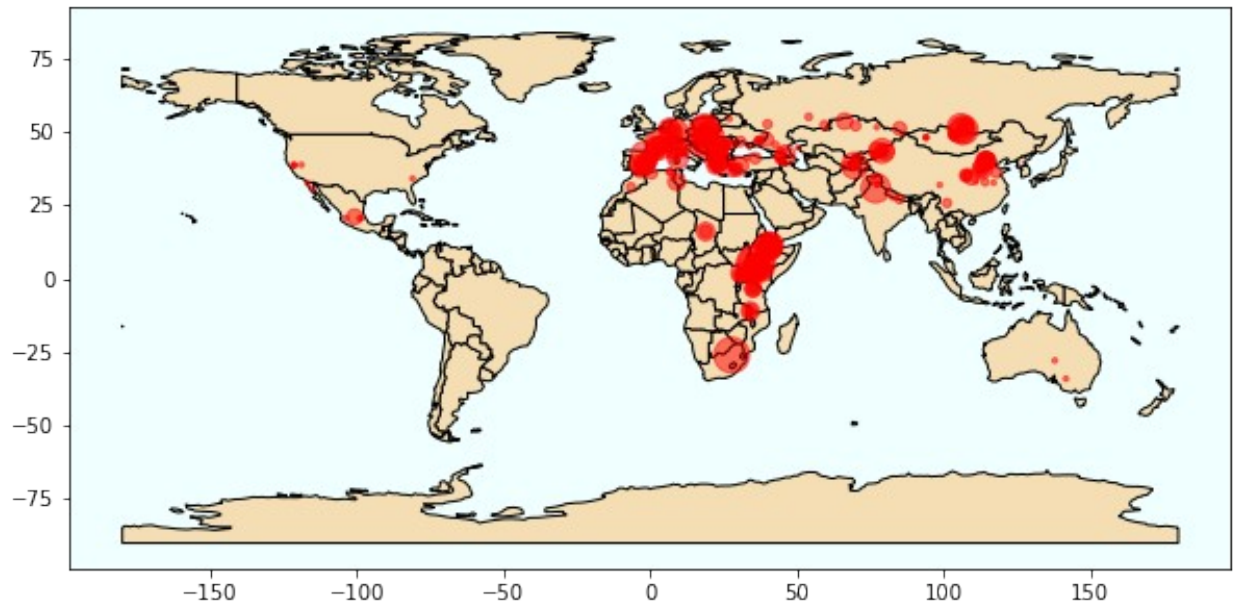
Time unit: MN14. Localities: 104



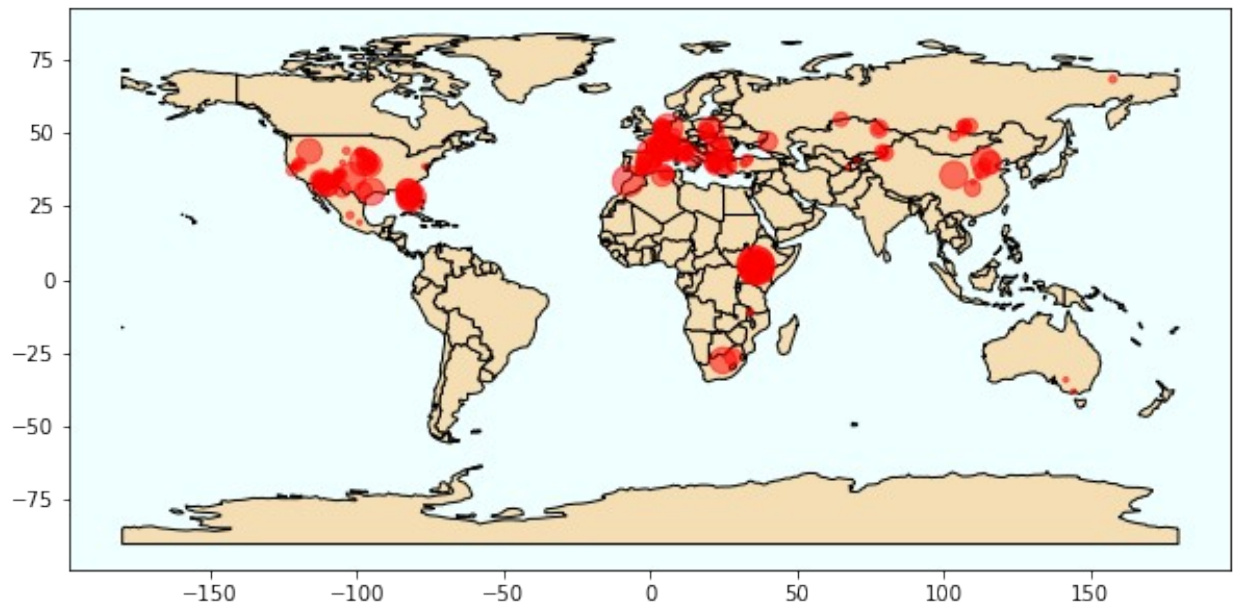
Time unit: MN15. Localities: 351



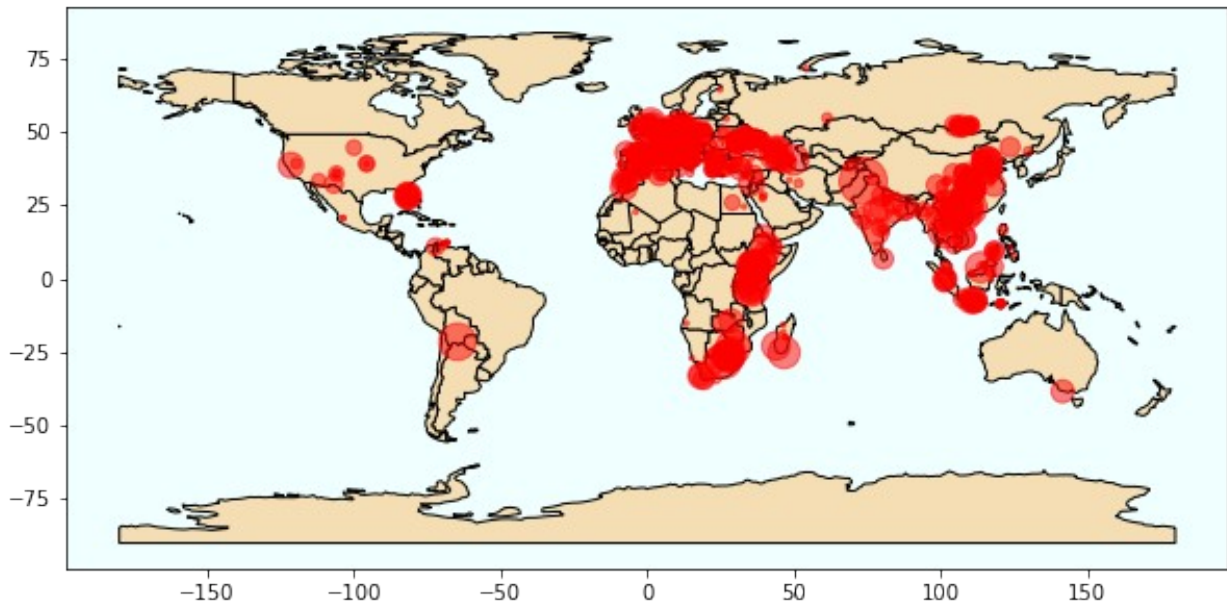
Time unit: MN16. Localities: 266



Time unit: MN17. Localities: 173



Time unit: post-MN. Localities: 659



I assumed the instructions meant that I will need to plot all of the 19 or so time units. So there they are. You can see how the occurrences seemed to be moving in waves; less occurrences -> more occurrences in the next time period -> ... and so on.

```
# Plotting localities for different continents based on instruction 5
C.
def plot_continental_info():
    locations = location_df()
    ax = world.plot(color='wheat', edgecolor='black', figsize=(20,
10))
    ax.set_facecolor('azure')

    # Creating geopandas DataFrame for joining
    gdf = geopandas.GeoDataFrame(locations,
geometry=geopandas.points_from_xy(locations.LONG, locations.LAT))
    gdf.set_crs(epsg=4326, inplace=True)

    # Joining locality data with the world map and grouping by
continent
    combined = geopandas.tools.sjoin(gdf, world,
how='left').groupby('continent')
    combined.plot(ax=ax, markersize=locations.ALL_OCC*5, alpha=0.5) #
Plotting the data with dynamic markersize

    continent_data = combined.agg('mean') # Getting average occurrences
for every locality in each continent

    plt.suptitle('Average number of occurrences per locality in every
```

```
continent.')
```

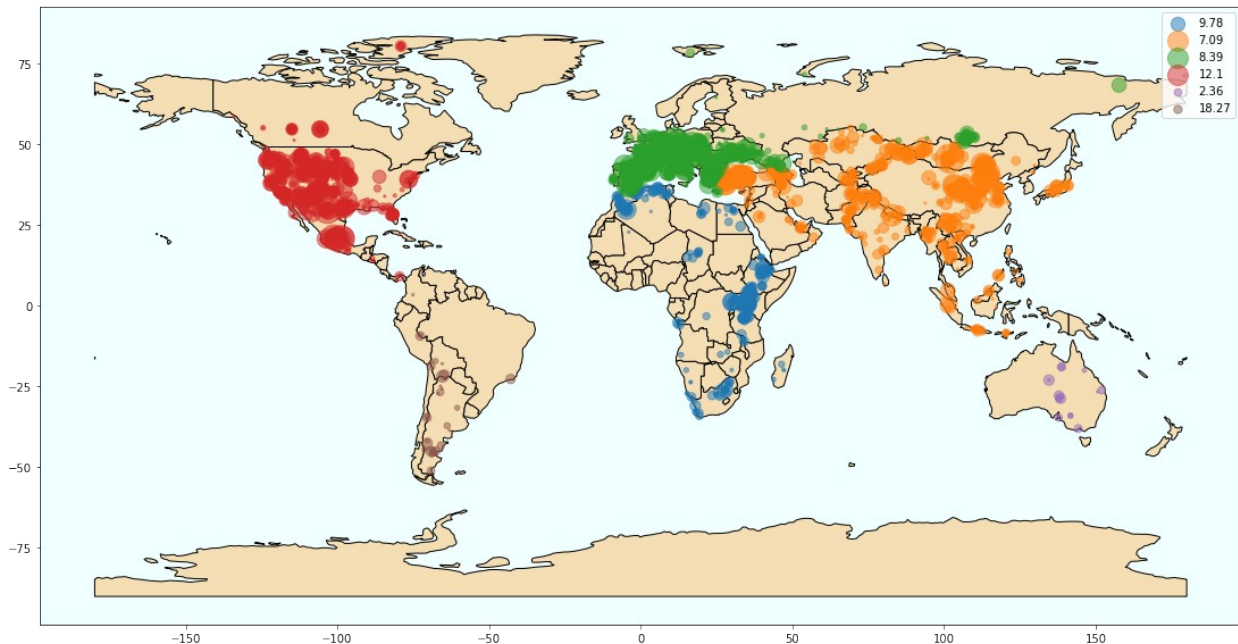
```
    ax.legend(round(continent_data.ALL_OCC, 2))
```

```
    plt.show()
```

```
if __name__ == '__main__':
```

```
    plot_continental_info()
```

Average number of occurrences per locality in every continent.



Answer

I should be using **LIDNUM** because some of the findings have the same locality names with different **LIDNUM**; it helps to further improve the accuracy of the analysis. (there also might be same names for different places)

1. c) There are some sampling density changes, most notable are **MN1**, **MN2** and **MN14**; **post-** and **pre-MN** aren't included since their sample size is much larger compared to the other time units. The rest is around the median. The second plot section shows average number of occurrences per locality for every continent. **South America** with the most occurrences (**18.27**) and **Australia** with the least amount of occurrences (**2.36**).

Localities and sampling

1. For every locality get 10x10 area of nearby localities and sum their all and first occurrences for the focal locality's time unit. Also get the previous time unit's all occurrences.

```

def check_previous_mn(idx, mn, mn_groups):
    tu = list(MN_UNITS.keys())
    previous_mn = tu.index(mn_groups[idx])-1
    if tu[previous_mn] == mn:
        return True
    return False

def square_localities():
    locations = location_df()
    # Creating geopandas DataFrame for joining
    gdf = geopandas.GeoDataFrame(locations,
    geometry=geopandas.points_from_xy(locations.LONG, locations.LAT))
    gdf.set_crs(epsg=3395, inplace=True)
    mn_groups = gdf.MN.to_dict() # Current time periods for every
    group in a dict is faster to use

    buffer = gdf.buffer(5, cap_style=3) # Create a 10x10 square around
    the localities
    buffer = geopandas.GeoDataFrame(geometry=buffer).to_crs(3395) #
    GeoSeries -> GeoDataFrame for faster intersection

    # Join and groupby index to get every 10x10 square of localities
    inside_area = geopandas.tools.sjoin(buffer, gdf, how='left',
    op='contains')
    inside_area[['FIRST', 'ALL']] =
    inside_area.groupby([inside_area.index, 'MN'])[['FIRST_OCC',
    'ALL_OCC']].transform(sum)

    # Getting all and first occurrences in the focal time unit
    focal_mn = inside_area[[True if mn_groups[idx] == mn else False
    for idx, mn in zip(inside_area.index, inside_area['MN'])]]
    focal_mn = focal_mn[~focal_mn.index.duplicated(keep='first')]

    # Getting all occurrences in the previous time unit
    before_focal = inside_area[[check_previous_mn(idx, mn, mn_groups)
    for idx, mn in zip(inside_area.index, inside_area['MN'])]]
    before_focal =
    before_focal[~before_focal.index.duplicated(keep='first')]

    # Adding new new data to the exercise 5 DataFrame
    locations = locations.assign(FOCAL_ALL=focal_mn.ALL,
    FOCAL_FIRST=focal_mn.FIRST, PREVIOUS_ALL=before_focal.ALL)
    locations.drop('geometry', axis=1, inplace=True) # Drop geometry
    column because it's not needed
    # locations.dropna(subset = ["PREVIOUS_ALL"], inplace=True) # Used
    to drop all rows without a previous time unit
    return locations

```

```
if __name__ == '__main__':  
    square_localities()
```

Answer

This probably isn't the optimal solution since it takes some time to run but It gives the correct `DataFrame`. The most time consuming part is the one instance of `geopandas sjoin` method with its staggering **2.2 million rows**. Giving `sjoin` the parameter `op='contains'` speeds it up a bit.

Notes: It wasn't specified what you should do if locality doesn't have a previous time unit so I returned NaN values for those since it seemed to be the right thing to do considering we are going to use the data for regression.